

Assignment 2

This assignment is based on an assignment given in the Machine Learning and Pattern Recognition course in our university. The first exercise is supposed to test you on your ability to load and manipulate data in python. This is an absolutely crucial skill you need to have to work on any kind of machine learning project. The second exercise might be a bit hard depending on your background. Linear regression is a rather intuitive method that can be effectively used on its own for some problems or as a baseline for more complex problems. You can use any kind of tutorials and materials online but please do not copy code from anywhere. The crucial part here is for you to actually understand how linear regression works so do not use ready made models but design the model yourself. You can however use an already made function for fitting the weights (`np.linalg.lstsq()`)

In terms of libraries use only numpy, any image processing library (eg matplotlib) and any kind of library you want for loading the data.

The data will be available on our GitHub.

1.

a) Plot a line graph showing the sequence in `amp_data`, and a histogram of the amplitudes in this sequence. Include these plots in your report, with one to three sentences about anything you notice that might be important for modelling these data.

b) Randomly shuffle the rows of the matrix. Then split the data into training (70%), validation (15%), and testing (15%). Each dataset should take the first $D = 20$ columns as a matrix of inputs X , and take the final column to create a vector of targets y . Name the resulting six arrays: `X_shuf_train`, `y_shuf_train`, `X_shuf_val`, `y_shuf_val`, `X_shuf_test` and `y_shuf_test`. The shuffling means that our training, validation and testing datasets all come from the same distribution. Include your code for creating the six arrays above from the original `amp_data` array.

2.

Given just one row of inputs, we could fit a curve of amplitude against time through the 20 points, and extrapolate it one step into the future. Plot the points in one row of your `X_shuf_train` data against the numbers $t = 0/20, 1/20, 2/20, \dots, 19/20$ representing times. We can fit this sequence with various linear regression models, and extrapolate them to predict the 21st time step at time $20/20 = 1$. Indicate the point you're predicting from `y_shuf_train` on the plot at $t=1$.

Fit and plot a straight line to the 20 training points using linear regression.